

STEFAN PEŠOVIĆ

Full-Stack Developer | AI Engineer | Node.js + React | Python + LLMs

Serbia / Czech Republic • stefan.pesovic13@gmail.com • github.com/stefanpesovic • linkedin.com/in/stefan-pesovic-389a893a3

ABOUT ME

Self-taught Full-Stack Developer and AI Engineer with 3+ years of Python and JavaScript experience. I build across the entire stack — Node.js/Express backends, React frontends, REST APIs — and integrate production-ready AI systems: RAG pipelines, LLM agents, and content automation. I don't just integrate APIs — I ship architected, tested, deployed products end-to-end, both greenfield and within existing codebases. My recent portfolio includes 3 production projects with 274+ tests, one with live deployment.

I use Claude Code and modern AI-assisted development tools as my daily driver. This means I move faster than traditional developers while maintaining production quality — I architect the system, write acceptance criteria, and direct implementation, then review and test everything that ships. This workflow is what forward-looking teams are adopting in 2026.

Beyond client work, I run a local AI server built around an RTX 5080 GPU for privacy-focused agents, and I'm actively building commercial SaaS products (AI waiter platform for restaurants, LingoDish — multilingual recipe platform for foreign kitchen staff).

FEATURED PORTFOLIO PROJECTS

VideoBrief — AI-Powered YouTube Summarizer with LangGraph Agent

Full-Stack Developer

Paste any YouTube URL → get a structured markdown report with key takeaways, clickable timestamps, and video-type-specific insights. Built with a 5-stage LangGraph agent that autonomously decides which specialized tools to run based on video content (tutorials trigger code extraction, interviews trigger quote extraction, etc.).

Tech Stack: Python, FastAPI, LangGraph, Groq (Llama 3.3 70B), youtube-transcript-api, React 18, Vite, TailwindCSS, Framer Motion, Docker

Key Achievements:

- Multi-stage LangGraph agent with conditional routing based on content type
- 137 tests (63 backend + 74 frontend), full CI/CD with GitHub Actions
- Glassmorphism UI with Hero, Feature Cards, Video Preview, and Staged Loader
- Adaptive reports: tutorial vs. interview vs. educational content get different structures

Repository: github.com/stefanpesovic/VideoBrief

AskStefan — Multi-Document RAG Chatbot (LIVE DEPLOYMENT)

Full-Stack Developer

Production RAG (Retrieval-Augmented Generation) chatbot that answers questions about me based on my CV, project descriptions, and personal documents. Features a glassmorphism split-layout UI with highlighted source citations, cold-start loader, and dismissible welcome banner.

Tech Stack: Python, FastAPI, LangChain, ChromaDB (vector DB), Cohere (embeddings), Groq (Llama 3.3 70B), React, TailwindCSS, Docker, Render

Key Achievements:

- LIVE demo: askstefan.onrender.com — anyone can interact with it
- Multi-document ingestion (PDF + Markdown) with SHA-256 hash deduplication

- 88 tests (67 backend + 21 frontend), Docker Compose, GitHub Actions CI
- Citation-aware retrieval: every answer shows exact source chunks with highlighted keywords

Live Demo: askstefan.onrender.com

Repository: github.com/stefanpesovic/AskStefan

AIJobRadar — Real-Time AI/ML Job Scraper REST API

Backend Developer

Production-grade REST API that scrapes real-time AI/ML job listings from RemoteOK, WeWorkRemotely, and Hacker News, normalizes them into a unified schema, and serves them via FastAPI. Async scraping completes all 3 sites in 1.4 seconds.

Tech Stack: Python 3.11, FastAPI, httpx (async), BeautifulSoup, Pydantic v2, pytest, Docker, GitHub Actions

Key Achievements:

- 72 tests, 89% code coverage
- Word-boundary regex filter solving substring false-positive bug (201 raw → 39 genuinely AI-relevant jobs)
- Async scraping with proper rate limiting and graceful error handling
- JSON cache with TTL for efficient re-querying

Repository: github.com/stefanpesovic/AIJobRadar

COMMERCIAL & PRIVATE PROJECTS

The following projects are in private repositories because they are commercial products or under active development for paying clients. Full code access available upon request via direct contact.

Best Guest — AI Waiter Platform for Restaurants (Active Development)

Restaurant platform where guests scan a QR code at their table and interact with an AI agent acting as a waiter. The agent knows the menu, ingredients, allergens, recommends dishes, and processes orders and payments directly through chat.

Tech Stack: Python, FastAPI, React, AI agents, payment integration, real-time chat

LingoDish — Multilingual Recipe SaaS (Live / Active)

SaaS platform solving a real hospitality problem: foreign kitchen staff who don't know local dish preparation standards. Restaurants subscribe and get a digital recipe book with ingredients, plating photos, and step-by-step preparation guides — translated into 12+ languages. Features a clean kiosk interface optimized for kitchen screens. Setup handled by our team within 24 hours. Currently live with 50+ active restaurants globally.

Tech Stack: TypeScript, React, backend API, multi-language translation, kiosk UI, 12+ languages

Live: lingodish.com

Email Warmup — Local Gmail Warmup System (Private)

Self-hosted tool for warming up Gmail accounts, running fully locally on Windows with no cloud services or third-party SaaS. Added Gmail accounts exchange naturally-looking emails on a daily schedule to build sender reputation. Features a password-protected web dashboard, support for unlimited accounts via App Passwords (no OAuth), 1–20 emails per account daily with 2–8 minute random delays, automatic replies and read-marking, rotating subject lines and message bodies (editable via Settings), warmup strength score 0–100% with 30-day ramp, activity log with pagination and CSV export, AES-256 encryption of credentials in local SQLite, and Windows Task Scheduler integration for automatic start/stop (10:00–22:00).

Tech Stack: Node.js, Express, SQLite, Nodemailer (SMTP), IMAP, node-cron, AES-256 encryption, vanilla JS/CSS frontend, Windows Task Scheduler

Apex Fitness — Personal Training Platform (Delivered to Client)

Web application where users purchase customized training and nutrition plans, receiving daily workouts and meals adapted to their goals. Project delivered and sold to a private client; currently in production.

Tech Stack: TypeScript, React, backend API, payment integration

Local AI Server — RTX 5080 Infrastructure

Built a dedicated AI server around an RTX 5080 GPU for running local models with full data privacy and zero recurring API costs. Hosts multiple custom agents: marketing content generator (Instagram, TikTok), coding assistant, image generation without cloud dependencies, and task-specific agents.

Tech Stack: Local LLM deployment, CUDA, Python agents, custom integrations

TECHNICAL SKILLS

Languages: Python (3+ years, daily), JavaScript/TypeScript (daily), Java (academic)

AI & LLMs: Anthropic Claude, OpenAI, Groq (Llama 3.3 70B), Cohere, LangChain, LangGraph, RAG systems, AI agents, prompt engineering, tool calling, vector databases

Backend: FastAPI (daily), async Python, Pydantic v2, REST API design, dependency injection, middleware, pytest, pytest-asyncio

Frontend: React 18, Vite, TailwindCSS, Framer Motion, Lucide React, Angular (working knowledge), responsive design, glassmorphism UI

Databases: MongoDB (daily), Mongoose, ChromaDB (vector), SQLite, basic SQL

DevOps & Deployment: Docker (multi-stage), docker-compose, GitHub Actions (CI/CD), Render platform, environment management, .env security

AI Development Tools: Claude Code (daily driver), Cursor, ChatGPT, local models on own GPU server

Rapid Prototyping: Streamlit, Gradio, FastAPI + Swagger UI

Methodology: PSM I certified (Professional Scrum Master I, Scrum.org), architect-first development, test-driven mindset, agile shipping

Operating Systems: Linux — Ubuntu (daily), macOS, Windows

EXPERIENCE

Freelance Full-Stack & AI Developer — *Self-Employed*

2023 – Present (3+ years)

Building AI-powered applications and automation systems for personal projects and clients. Focus on LLM integration, RAG systems, AI agents, production-grade backends with FastAPI and Node.js, and content production pipelines. Comfortable owning features end-to-end — from concept to shipped — and navigating existing codebases to deliver reliable outcomes.

- 3 production AI projects with 274+ tests across full stack
- Delivered commercial fitness SaaS (TypeScript + React) to private client
- Active development on AI waiter platform and LingoDish (multilingual recipe SaaS, 50+ restaurants live)
- Daily work with Claude Code for architecture, implementation, and testing; background jobs and scheduled automation (node-cron, Task Scheduler)

EDUCATION & CERTIFICATIONS

Master's in Economics — *University MB, Belgrade*

2025 – Present

Bachelor's in Economics — *Union University 'Nikola Tesla', Faculty of Management*
2021 – 2025

Professional Scrum Master I (PSM I) — *Scrum.org*

Python (course) • **Java (course)** • **ChatGPT & AI (specialized course)**

LANGUAGES

English: Professional working proficiency (B2-C1)

Czech: Fluent

German: Conversational

Serbian / Croatian / Bosnian: Native